

# SYZOS: practical KVM fuzzing

Alexander Potapenko <glider@google.com>  
Linux Plumbers Conference 2025

# The Challenge of Fuzzing KVM

1. Host (L0): complex, non-atomic setup:
  - Requires valid sequences of KVM\_SET\_CPUID2, KVM\_SET\_SREGS, KVM\_SET\_USER\_MEMORY\_REGION etc.
2. Guest (L1/L2): existing solutions within syzkaller were limited:
  - Strategy A: Hardcoded Assembly Blobs
    - Injecting static blobs via legacy setup calls.
    - Limitation: Zero flexibility. The fuzzer cannot mutate parameters or logic, restricting coverage to a single "happy path."
  - Strategy B: Pseudo-Instructions
    - Generating macros (e.g., PSEUDO\_WRMSR, PSEUDO\_HYPERCALL).
    - Limitation: These lower to raw assembly. The fuzzer mutates individual bytes, destroying the logic immediately.
3. Reachability: Some CPU state is inaccessible to L0 ioctls.
4. Coordination: Complex scenarios require turn-based execution (e. g. L0 ↔ L1).

# Foundation: Fix KVM Descriptions

- Not too many updates of `sys/linux/dev_kvm.txt` since 2019
- 50+ changes refining the descriptions to prune the search space
  - Opaque blobs → architecture-specific structs
  - Abstract int64 → constants or flags
  - New resource types to properly chain the syscalls
  - Splitting polymorphic ioctls (e.g. `KVM_GET_MSRS` is both a system and a VCPU ioctl)
  - Apply `no_squash` for complex struct argument to forbid arbitrary bitflips

# Host: pseudo-syscalls to the rescue

- `syz_kvm_setup_syzos_vm()`:
  - `ioctl(vmfd, KVM_SET_USER_MEMORY_REGION, &memreg) x N`
- `syz_kvm_add_vcpu$x86()`:
  - `ioctl(vmfd, KVM_CREATE_VCPU, cpu_id)`
  - `ioctl(cpufd, KVM_GET_SREGS, &sregs)`
  - Set up segments, page table, IDT
  - `ioctl(cpufd, KVM_SET_SREGS, &sregs)`

Syzkaller will generate these syscalls along with the normal ones, increasing the chance to hit more interesting paths.

They are implemented in C and are part of `syz-executor`.

# L1 Guest: Introducing SYZOS

SYZOS is an immutable C library with an easy-to-fuzz API.

The fuzzer generates a sequence of API calls (commands) for L1 to execute.

SYZOS primitives for x86 (ARM is also supported):

- Exit, execution, I/O:
  - SYZOS\_API\_UEXIT – trigger a page fault, return 1 argument to Host.
  - SYZOS\_API\_CODE – execute a Host-supplied instruction blob.
  - SYZOS\_API\_IN\_DX, SYZOS\_API\_OUT\_DX – execute the in and out instruction (port in DX).
- Privileged operations:
  - SYZOS\_API\_CPUID – execute the cpuid instruction.
  - SYZOS\_API\_WRMSR, SYZOS\_API\_RDMSR – read/write MSRs.
  - SYZOS\_API\_WR\_CRN, SYZOS\_API\_WR\_DRN – write CRO..CR8, DRO..DR7.
- Interrupt handling:
  - SYZOS\_API\_SET\_IRQ\_HANDLER – install an IRQ handler for the given interrupt.

# Nested Virtualization: x86 L2 Primitives

SYZOS acts as a lightweight L1 hypervisor to fuzz nested virtualization. It abstracts architectural differences (VMX vs SVM) into high-level commands.

SYZOS Primitives:

- VM Lifecycle Setup:
  - SYZOS\_API\_ENABLE\_NESTED – Enable virtualization extensions (VMX/SVM).
  - SYZOS\_API\_NESTED\_CREATE\_VM – Initialize VMCS/VMCB and Nested Page Tables.
  - SYZOS\_API\_NESTED\_LOAD\_CODE – Inject small instruction blobs into L2 memory.  
(Note: Currently supports code snippets, not the full SYZOS kernel)
- Execution Control:
  - SYZOS\_API\_NESTED\_VMLAUNCH – Perform initial entry into the L2 Guest.
  - SYZOS\_API\_NESTED\_VMRESUME – Re-enter the L2 Guest after an exit.
- State Mutation:
  - SYZOS\_API\_NESTED\_\*WRITE\_MASK – Mutate VMCS/VMCB control fields.

# ARM64 Specifics

## 1. Fuzzing the GICv3 ITS (Interrupt Translation Service)

- The ITS relies on complex in-memory command queues and translation tables.
- SYZOS Primitives:
  - SYZOS\_API\_ITS\_SETUP: Allocates tables and configures the GIC base.
  - SYZOS\_API\_ITS\_SEND\_CMD: Injects structured GIC commands (e.g., MAPD, MOVI) into the queue to fuzz the kernel's command processor.

## 2. Fuzzing the Hypervisor Interface (LO Traps)

- Testing the boundary between the Guest and the LO Hypervisor/Firmware.
- Hypercalls: SYZOS\_API\_HVC / SYZOS\_API\_SMC with fuzzer-controlled registers.
- Goal: Fuzz LO handlers for PSCI (Power State Coordination) and SMCCC (SiP service) calls.

# Execution Model

- Finite execution:
  - syzkaller generates a finite, end-to-end program (sequence of system calls).
  - The "Run Loop" is unrolled into discrete KVM\_RUN syscalls.
- The yield mechanism:
  - SYZOS\_API\_UEXIT is the explicit handshake to yield control from L1 to L0.
  - This allows the fuzzer to interleave L0 mutations (e.g., memory unmapping) precisely between Guest steps.
- Concurrency:
  - Syzkaller can mark syscalls as async to provoke Host/Guest data races.


# A Case Study

```
# Setup (not shown on the diagram):
```

```
r0 = openat$kvm(...)
r1 = ioctl$KVM_CREATE_VM(r0, ...)
r2 = syz_kvm_setup_syzos_vm$x86(...)
r3 = syz_kvm_add_vcpu$x86(
    ...,
    [@enable_nested,
     @nested_create_vm,
     @nested_load_code("f40f08"),
     @nested_vmlaunch,
     @nested_intel_vmwrite_mask(
         {0x4002, 0x0, 0x80, 0x0}
     ),
     @nested_vmresume
    ])
)
```

```
# Start:
```

```
ioctl$KVM_RUN(r3)
```




# A Case Study (cont.)

```
# Continue after UEXIT:
```

```
ioctl$KVM_RUN(r3)
```

```
# Continue after HLT:
```

```
ioctl$KVM_RUN(r3)
```



# Validation and Findings

## Qualitative impact

- Early fruit: **9 bugs** with SYZOS-based reproducers reported in 2025.
- "Collateral Damage": overhauling KVM syscall descriptions as part of SYZOS development.
- Deeper state access: replaced rigid "Happy path" instruction blobs with mutable logic to reach previously invisible states.

## Corpus adoption

- The fuzzer actively chooses SYZOS programs (20–30% of the programs in the minimal x86 KVM corpus, 3–4x growth in the past two months).
- 294 Intel and 157 AMD programs successfully trigger nested VM launches.

# Coverage Improvements

Fuzzing on GCE ([syzbot.org](https://syzbot.org)): 200+ machines, different kernel versions.

- arch/arm64/kvm coverage\* up by 5% with SYZOS rollout.
- x86-related KVM coverage up by 7–12%.

Fuzzing on bare-metal: 1–3 machines per architecture, single kernel, disabled the legacy guest blobs.

- Intel and AMD-specific KVM coverage doubled since SYZOS adoption (Intel: 58%, AMD: 54%).
- Legacy guest blobs were disabled to prove that the coverage gains are purely driven by SYZOS logic, not the old "happy path" instructions.

\* Here and below we mean line coverage that is easier to aggregate. No 1:1 correspondence to basic block coverage.

# Future Plans

- Deeper nesting: run SYZOS inside SYZOS.
- Better ARM support: page tables, NV, devices.
- Non-64-bit guest modes on x86.
- Analyze coverage gaps.
- Continuous testing.